```
import nltk
import numpy as np
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics.pairwise import cosine_similarity

# Load FAQs
faq_pairs = []
with open('faqs.txt', 'r') as f:
    lines = f.read().strip().split('\n')
    for i in range(0, len(lines), 3):
        q = lines[i][3:]
        a = lines[i+1][3:]
        faq_pairs.append((q, a))

# Preprocess with TF-IDF
vectorizer = TfidfVectorizer()
questions = [q for q, _ in faq_pairs]
tfidf_matrix = vectorizer.fit_transform(questions)

def get_response(user_input):
    input_vec = vectorizer.transform([user_input])
    similarities = cosine_similarity(input_vec, tfidf_matrix)
    idx = np.argmax(similarities)
    return faq_pairs[idx][1]

# Simple CLI chatbot
print("Welcome to the FAQ Bot! Type 'exit' to quit.")
while True:
    user_input = input("You: ")
    if user_input.lower() == 'exit':
        break
    response = get_response(user_input)
    print("Bot:", response)
```